

marp: true

theme: gaia

size: 16:9

paginate: true

## Deep Q-Network Solution for Stochastic Gridworld Navigation

### A Reinforcement Learning Project

#### The Problem: A Stochastic Gridworld

The goal is to train an agent to find the best path in a 3x4 grid.

Goal: (0,3) | Reward: +1.0

Trap: (1,3) | Reward: -1.0

Obstacle: (1,1) | Impassable Wall

Step Cost: Every move costs -0.04 to encourage efficiency.